

Architectural Foundations of a Self-Improving Mixture of Experts for Financial Asset Forecasting

Forecasting equity indices and commodity futures represents one of the most demanding frontiers in applied machine learning, characterized by severely degraded signal-to-noise ratios, pronounced fat-tailed distributions, time-varying autoregressive dependencies, and pervasive structural breaks driven by shifting macroeconomic regimes. Conventional deep learning models trained under static objectives inevitably suffer from alpha decay and out-of-sample performance collapse as the underlying data-generating process drifts. While standard sparse Mixture of Experts (MoE) architectures provide scalable parameter capacity by conditionally routing input tokens across parallel feed-forward networks, standard routing formulations remain ill-suited for financial time series because their gating mechanisms are memoryless, assume stationary input distributions, and fail to optimize downstream financial utility. Engineering a self-improving MoE capable of generating sustained predictive performance across global indices and commodity complexes requires an end-to-end quantitative framework. This encompasses an information-matched routing backbone that decouples macroeconomic regime discovery from localized asset pricing signals, a multi-modal feature ingestion pipeline that respects microstructure invariants and commodity physical constraints, a continual learning engine trained via direct reinforcement learning on differential risk-adjusted reward functions, and a validation harness grounded in combinatorial purging and multiple-testing corrections.

Sparse Mixture of Experts Architectures for Non-Stationary Financial Dynamics

The Regime-Gated Information Decoupling Principle

A primary pathology in financial time-series modeling is the direct concatenation of non-stationary macroeconomic and market-state indicators with localized asset-level features. When wide-spectrum regime indicators—such as implied volatility surfaces, sovereign yield curves, credit default swap spreads, and systemic liquidity measures—are passed directly into the same predictive layers as asset-level returns and order flow, deep networks experience severe optimization instability and negative transfer. The low-frequency, high-variance dynamics of macro variables overwhelm the subtle, high-frequency predictive patterns embedded within asset-specific features.

The Regime-Gated Residual Mixture of Experts (RG-ResMoE) architecture resolves this pathology by establishing an information-matched integration pathway. The model partitions the input domain into a localized asset feature vector $\mathbf{x}_t \in \mathbb{R}^{d_x}$, containing stationary or fractionally differentiated time series, and an external market regime vector $\mathbf{z}_t \in \mathbb{R}^{d_z}$, encapsulating broad market volatility and liquidity states. Crucially, the predictive forecasting blocks are evaluated exclusively over $\mathbf{x}_t$, while the regime state $\mathbf{z}_t$ affects predictions solely through the gating parameterization:

$$\hat{y}_t = f_{\text{base}}(x_t; \theta_0) + \sum_{i=1}^E g_i(z_t; W_g) \cdot f_{\text{expert}, i}(x_t; \theta_i)$$

In this residual formulation, the primary baseline predictor f_{base} models unconditional asset drift and structural trend components, while E localized residual experts $f_{\text{expert}, i}$ learn state-dependent adjustments. Initializing the residual expert output projections near zero ensures that the architecture begins training as a stable base forecaster, gradually incorporating expert specialization as optimization progresses.

The routing gating function $g(z_t)$ employs continuous soft routing via a parameterized softmax operator:

$$g(z_t) = \text{Softmax}(W_g z_t + b_g)$$

Empirical evidence indicates that continuous soft routing strictly outperforms hard Top-k selection across out-of-sample financial evaluation windows. Financial regimes do not transition along discrete, deterministic step functions; rather, market states evolve across continuous latent distributions. Soft routing allows the model to blend specialized modules smoothly during intermediate, transitional market regimes—such as an orderly transition from low-volatility monetary accommodation into high-volatility inflation spikes—thereby preventing the abrupt parameter discontinuities and activation chatter characteristic of hard routing.

Variable-Context Expert Networks and Multi-Scale Patching

A persistent limitation of time-series backbones in quantitative finance is the enforcement of a uniform context window length across all temporal steps. Financial dynamics alternate between high-speed transactional volatility, where optimal lookback horizons are short, and extended macroeconomic trends, where extended context windows are mandatory to capture secular momentum. A static context window inevitably over-smooths local microstructure shifts if set too long, or truncates macroeconomic auto-dependencies if set too short.

The Regime-Aware Variable-context Expert Network (RAVEN) addresses this challenge by routing varying temporal lookback contexts to specialized sub-networks. By applying Cumulative Importance Thresholding (CIT) to the attention scores across historical input segments, the architecture establishes a nested sequence of lookback horizons $\tau_1 \subset \tau_2 \subset \dots \subset \tau_E$, where each horizon is processed by an expert specialized at that specific temporal scale.

Because nested temporal contexts share historical data points, unconstrained expert outputs exhibit severe multicollinearity and mutual information overlap. To enforce functional diversity and eliminate redundant representations across the ensemble, Correlation-Aware Weighting (CAW) applies a cosine penalization across the terminal hidden states h_i of the individual experts prior to output aggregation:

$$\hat{y}_t = \sum_{i=1}^E \tilde{w}_i(z_t) \cdot f_i(x_t^{\tau_i}) \quad \tilde{w}_i = \frac{g_i(z_t) \exp\left(-\lambda_c \sum_{j \neq i} \cos(h_i, h_j)\right)}{\sum_{k=1}^E g_k(z_t) \exp\left(-\lambda_c \sum_{j \neq k} \cos(h_k, h_j)\right)}$$

where $\lambda_c > 0$ governs the repulsion magnitude against collinear hidden vectors. This explicit regularization forces the router to maximize the marginal information gain of each additional lookback scale, optimizing ensemble diversity.

Sequence Modeling Backbones and Linear Scaling

For the internal expert blocks, Transformer-based self-attention suffers from quadratic computational complexity $\mathcal{O}(L^2)$ and high memory footprints when processing long

lookback sequences. In contrast, selective state-space architectures (such as Mamba and MambaFormer-MoE) achieve linear time complexity $\mathcal{O}(L)$ and constant-state memory consumption during inference by replacing non-local self-attention with input-dependent continuous recurrent projections. This linear scaling enables the ingestion of extended high-frequency tick histories and multi-scale macro series without truncating context. Concurrently, channel-independent patching frameworks—such as PatchTST and Time-MoE—segment univariate series into localized sub-series tokens. Tokenizing localized temporal patches rather than discrete, noisy scalar points enhances the signal-to-noise ratio by extracting localized morphological structure while insulating the network against high-frequency gaussian noise. Integrating selective state-space layers into the longer-context experts while utilizing compact patched feed-forward blocks in the short-horizon residual experts yields an optimal balance between sequence modeling capability and inference efficiency.

Architecture Type	Routing Mechanism	Information Pathway	Complexity per Step	Target Financial Failure Mode
Standard Sparse MoE	Top-k Hard Gating	Concatenated (x_t, z_t)	$\mathcal{O}(k \cdot L^2)$	Optimization instability; representation collapse
RG-ResMoE	Continuous Soft Softmax	Isolated ($z_t \mapsto \text{Gate}, x_t \mapsto \text{Experts}$)	$\mathcal{O}(E \cdot L_{\text{compact}})$	Negative transfer from macro distribution shift
RAVEN	Cumulative Importance Thresholding	Nested multi-scale contexts τ_i	$\mathcal{O}(\sum_{i=1}^E L_i^2)$	Fixed lookback mismatch to market volatility
MambaFormer-MoE	Dynamic Cross-Variate Gate	Selective State-Space Recurrence	$\mathcal{O}(E \cdot L)$	Memory exhaustion under long lookback contexts

Multi-Asset Feature Engineering and Structural Data Representation

Information-Preserving Transformation via Fractional Differentiation

The conventional methodology for enforcing stationarity in financial modeling relies on integer first-differencing ($d=1$) of raw prices. Although first-differencing establishes weak stationarity, it operates as an aggressive high-pass filter that eradicates the multi-period memory of the underlying price process, destroying long-range mean-reverting and trending properties. Conversely, preserving raw price levels ($d=0$) retains complete memory but introduces non-stationary distributions that violate econometric stability criteria and destabilize neural network gradient dynamics. Fractional differentiation achieves stationarity while preserving the maximum degree of historical memory. By expanding the binomial series of the difference operator $(1-B)^d$ for real-valued d : $(1-B)^d = \sum_{k=0}^{\infty} (-1)^k \binom{d}{k} B^k = 1 - d B + \frac{d(d-1)}{2!} B^2 - \frac{d(d-1)(d-2)}{3!} B^3 + \dots$ where B denotes the backshift operator ($B^k X_t = X_{t-k}$), the memory weights ω_k are evaluated recursively:

$$\omega_0 = 1, \quad \omega_k = -\omega_{k-1} \frac{d - k + 1}{k}, \quad k \geq 1$$

To deploy this transformation, the minimal stationarity order d^* is calculated independently for every asset index and commodity futures contract via an expanding Augmented Dickey-Fuller (ADF) grid search. The parameter d^* is established as the lowest real value $d \in [0, 1]$ that rejects the non-stationarity null hypothesis at the 95% confidence level ($p < 0.05$). For equity indices and liquid energy contracts, empirical d^* values typically settle between 0.35 and 0.55, successfully preserving significant memory of historical cycles while eliminating distribution drift in the feature space.

Microstructure Dynamics and High-Frequency Imbalance

At intra-day forecasting horizons, closing prices fail to capture the supply-demand dynamics operating within the limit order book. The Order Flow Imbalance (OFI) metric formalizes the net shift in liquidity provision and consumption across discrete order book states. Let P_t^b , q_t^b represent the highest active bid price and bid size at time t , and let P_t^a , q_t^a represent the lowest active ask price and ask size. The Order Flow Imbalance across consecutive updates is defined as:

$$\text{OFI}_t = I_{\{P_t^b \geq P_{t-1}^b\}} q_t^b - I_{\{P_t^b \leq P_{t-1}^b\}} q_{t-1}^b - I_{\{P_t^a \leq P_{t-1}^a\}} q_t^a + I_{\{P_t^a \geq P_{t-1}^a\}} q_{t-1}^a$$

where $I_{\{\cdot\}}$ is the indicator function.

To eliminate microstructure noise and bid-ask bounce, scalar OFI_t values are accumulated across volume-based bars rather than fixed calendar time intervals. Integrating cumulative OFI with cross-sectional depth variations provides a direct measurement of latent adverse selection pressure. This gives the short-horizon residual experts an empirical leading indicator of dealer inventory rebalancing prior to subsequent price movements.

Cross-Asset Signals: Equities versus Commodity Physical Fundamentals

Constructing a unified forecasting framework for equity indices and commodity futures requires distinct feature sets that reflect their underlying economics. While broad equity indices represent financialized claims on corporate cash flows influenced primarily by systemic discount rates and macro liquidity, commodities are physical assets subject to non-linear physical storage limits, transport constraints, and localized supply shocks.

Equity index regimes are governed by macroeconomic liquidity flows, corporate risk appetites, and sovereign term structures. Feature pipelines must capture systemic stress via credit default swap (CDS) spreads, sovereign yield curve slopes (the 10-year minus 2-year Treasury spread), monetary liquidity indices (such as the trade-weighted US Dollar index DXY), and volatility surface metrics. The spread between implied volatility and trailing realized volatility isolates the variance risk premium, signaling institutional hedging demand and equity crash risk.

Commodity markets require structural modeling of physical supply-demand dynamics, futures term structures, and commercial hedging pressures:

- Term Structure and Roll Yield Dynamics:** The slope of the commodity futures curve reflects inventory availability. The annualized roll yield between the front-month contract (F_1) and second-month contract (F_2), defined as $\ln(F_1 / F_2) \cdot (365 / \Delta t)$, quantifies whether a market is in backwardation ($F_1 > F_2$, reflecting inventory scarcity and high convenience yields) or contango ($F_1 < F_2$, reflecting inventory surplus and

- positive storage costs).
- **Speculative Imbalance and Positioning:** Data from the CFTC Commitments of Traders (COT) reports provides institutional positioning insight. The normalized non-commercial net positioning ratio, computed as speculative long contracts minus short contracts divided by aggregate open interest, isolates overextended speculative positioning susceptible to rapid liquidation cascades.
 - **Physical Inventories and Seasonal Baselines:** Periodic balance sheets, such as the USDA World Agricultural Supply and Demand Estimates (WASDE) and the EIA Petroleum Status Report, introduce structural price shifts. Inventory metrics must be seasonally adjusted by standardizing current volumes against trailing five-year historical averages to isolate unpriced supply deficits or surpluses from predictable seasonal cycles.

Event-Driven Sampling and Triple-Barrier Meta-Labeling

Financial time series sampled at constant calendar frequencies (such as daily or hourly closes) exhibit high heteroskedasticity and non-Gaussian returns. Periods of market inactivity generate redundant, uninformative data points, while periods of market turbulence compress substantial trading volume and volatility into single time steps.

To align model updates with the true rate of information arrival, the architecture uses a symmetric Cumulative Sum (CUSUM) filter to drive event-based sampling. Given an asset's continuous return process r_t , the CUSUM detector accumulates directional return innovations against an adaptive volatility baseline σ_t :

$$S_t^+ = \max(0, S_{t-1}^+ + r_t - E[r_t]), \quad S_t^- = \min(0, S_{t-1}^- + r_t - E[r_t])$$

An observation event is triggered at timestamp τ when $\max(S_\tau^+, -S_\tau^-) \geq h \cdot \sigma_\tau$, at which point the cumulative metric is reset to zero. This aligns feature extraction with market information flow rather than arbitrary clock time.

Labels are generated using the Triple-Barrier Method, replacing fixed-time horizon returns. At each event timestamp t_0 , three dynamic thresholds are created: an upper horizontal barrier $P_0(1 + k \sigma_t)$ for profit taking, a lower horizontal barrier $P_0(1 - k \sigma_t)$ for stop-loss mitigation, and a vertical barrier $t_0 + \Delta t$ defining maximum trade duration. The primary directional target $y \in \{-1, \text{start_span} \text{end_span} + 1\}$ is determined by whichever horizontal barrier is breached first.

To optimize trade sizing, the framework separates directional classification from bet sizing through meta-labeling. The base MoE generates a primary directional forecast $\hat{y}_t \in \{-1, +1\}$. A secondary model estimates the probability that the trade will hit the profit-taking threshold before hitting the stop-loss barrier or expiring vertically. This secondary probability p_t is mapped to a continuous trade sizing coefficient $s_t \in [0, 1]$ via a calibrated sigmoid transformation:

$$s_t = \max\left(0, 2 \cdot \left(\sigma \left(\frac{p_t - 0.5}{\sigma_{\text{cal}}}\right) - 0.5\right)\right)$$

This two-stage formulation decouples the directional signal from execution conviction, enabling the strategy to downscale or eliminate position allocations during periods of market ambiguity.

Feature Class	Raw Input Sources	Asset Coverage	Mathematical Representation	Economic Function
Microstructure Flow	LOB Depth (Best 5 Ticks), Direct	Index & Energy Futures	Cumulative OFI_t across	Captures institutional order

Feature Class	Raw Input Sources	Asset Coverage	Mathematical Representation	Economic Function
	Feeds		volume bars	pressure and dealer imbalances
Fractional Memory	Unadjusted Continuous Settlement Prices	All Commodities & Indices	$(1-B)^{d^*} X_t$, d^* via ADF test	Preserves long-range memory while satisfying stationarity
Term Structure	WTI, Brent, Natural Gas, Corn Curves	Physical Commodity Futures	$\frac{1}{n} \sum_{t=0}^{n-1} (F_1 / F_2) \cdot (365 / \Delta t)$	Measures convenience yields and inventory storage pressure
Speculative Pressure	CFTC Commitments of Traders (COT)	Equity Futures & Ags	$\frac{\text{NonCommitted}_{Long} - \text{NonCommitted}_{Short}}{\text{OI}_{Total}}$	Identifies crowded positioning and short-squeeze risks
Variance Premium	VIX, OVX, GVZ, Trailing Realized Vol	Indices, Crude, Gold	$\text{ImpliedVol} - \text{RealizedVol}$	Measures systemic hedging demand and market crash risk

Continual Learning and Dynamic Self-Improvement Frameworks

Direct Reinforcement Learning with Differential Risk Ratios

Traditional supervised learning objectives minimize symmetric scalar errors such as Mean Squared Error (MSE), implicitly treating upside and downside forecast deviations identically. In live trading, however, a minor forecast error during calm drift regimes has negligible financial impact, whereas an identical absolute error during a severe liquidity shock can lead to extreme drawdowns.

To align optimization with trading performance, the self-improving gating policy is trained via reinforcement learning directly maximizing the Differential Sharpe Ratio (DSR). The sample Sharpe ratio cannot be computed recursively in a streaming online environment without re-evaluating the entire historical trajectory at every step. The Moody-Saffell Differential Sharpe

Ratio resolves this by expanding the Sharpe ratio to first order around an exponential moving average decay parameter $\eta \in (0, 1)$:

$$D_t \equiv \frac{\partial S_t}{\partial \eta} = \frac{B_{t-1} \Delta A_t - \frac{1}{2} A_{t-1} \Delta B_t}{(B_{t-1} - A_{t-1})^2}^{3/2}$$

where A_t and B_t represent recursive exponential moving averages of the first and second non-central moments of realized portfolio returns R_t :

$$A_t = A_{t-1} + \eta (R_t - A_{t-1}) = A_{t-1} + \eta \Delta A_t \quad B_t = B_{t-1} + \eta (R_t^2 - B_{t-1}) = B_{t-1} + \eta \Delta B_t$$

The realized trading system return R_t explicitly accounts for leverage, directional positioning F_{t-1} , and variable transaction costs δ :

$$R_t = F_{t-1} r_t - \delta |F_t - F_{t-1}|$$

To penalize downside risk more heavily, the objective function is extended to the Differential Downside Deviation Ratio (D3R), directly optimizing the Sortino ratio online:

$$D_t^{\text{Sortino}} = \begin{cases} \frac{R_t - \frac{1}{2} A_{t-1}}{DD_{t-1}}, & R_t > 0 \\ \frac{DD_{t-1}^2 (R_t - \frac{1}{2} A_{t-1}) - \frac{1}{2} A_{t-1} R_t^2}{DD_{t-1}^3}, & R_t \leq 0 \end{cases} \quad DD_t^2 = DD_{t-1}^2 + \eta (\min(R_t, 0)^2 - DD_{t-1}^2)$$

Optimizing the routing weights W_g directly against D_t^{Sortino} via policy gradient methods penalizes portfolio return volatility driven by downside tail risk. The gating network dynamically allocates capital toward defensive, low-beta experts whenever market conditions indicate accelerating downside volatility.

Preventing Catastrophic Forgetting via Modular Policy Libraries

Continual online fine-tuning of neural network parameters causes catastrophic forgetting, where updates learned in newly encountered environments overwrite previously acquired representations. In quantitative trading, an MoE model fine-tuned exclusively during an extended multi-year equity bull market overwrites the parameters of its crash-protection modules, leaving the system vulnerable during subsequent liquidity shocks.

Standard continual learning remedies, such as Elastic Weight Consolidation (EWC) and Experience Replay (ER), show limited effectiveness when applied to non-stationary financial time series. EWC's quadratic parameter penalties assume a stationary loss basin, leading to underfitting in shifting financial regimes, while Experience Replay suffers from severe class imbalance and recency bias.

The Regime-aware Continual Adaptive Portfolio (ReCAP) framework resolves this vulnerability through modular parameter isolation. The architecture maintains a frozen base parameter vector θ_0 trained across broad historical cycles. When an Adaptive Regime Detection (ARD) module—driven by real-time CUSUM change-point statistics on market features—detects a structural regime shift, the architecture isolates and trains a modular policy vector:

$$d_k = \theta_k - \theta_0$$

These specialized vectors are consolidated into an expanding policy library $\mathcal{D} = \{d_1, d_2,$

$\dots, d_K\}$. During live inference, when a new market state is detected, the base parameters remain unchanged while a dedicated delta vector d_{new} is instantiated and optimized on the incoming data stream.

Simultaneously, a soft gating module dynamically computes allocation weights $w_k(z_t)$ over the existing policy library, synthesizing the active network parameters dynamically:

$$\theta_{\text{active}} = \theta_0 + \sum_{k=1}^K w_k(z_t) \odot d_k$$

This modular formulation eliminates catastrophic forgetting. Established expertise remains preserved in the policy library, enabling the gating mechanism to instantly reactivate historically validated defensive parameters the moment crisis-like macro conditions reappear.

Counterfactual Credit Assignment and Retrospective Reflection

A persistent limitation of standard MoE routing is that routers receive no direct gradient feedback regarding unselected, inactive experts. If an expert is unselected by the gating network, its parameters receive zero updates, leaving the router unaware of cases where that inactive expert would have outperformed the active ensemble.

To enable genuine self-improvement, the gating mechanism incorporates counterfactual expert credit assignment. For each completed trading decision window t , the system samples inactive experts and calculates their counterfactual marginal loss reduction. Using a first-order Taylor expansion over the downstream task loss $\mathcal{L}$, the counterfactual credit $C_{i,t}$ for replacing active expert a with inactive alternative i is calculated as:

$$C_{i,t} \approx \mathcal{L}(\hat{y}_t^{(a)}) - \mathcal{L}(\hat{y}_t^{(i)}) \approx -\nabla_{\hat{y}_t} \mathcal{L}(\hat{y}_t)^{\top} \left(f_i(x_t) - f_a(x_t) \right)$$

If $C_{i,t} > 0$, the router's query-key memory matrices are updated using a margin-based contrastive loss, ensuring that unselected experts that would have improved out-of-sample performance receive higher routing probability under similar market conditions in the future. Simultaneously, an episodic retrospective reflection loop runs post-settlement. Daily performance records—encompassing realized returns, transaction costs, execution slippage, and regime classification confidences—are saved to an episodic memory buffer. When realized losses exceed a predefined risk threshold (such as three trailing standard deviations of forecast error), an automated reflection engine analyzes the prediction trace. The system adjusts the router's sensitivity parameters to downweight overconfident experts during transitional market regimes, preventing recurring prediction errors in similar macro conditions.

Objective Functions and Optimization in Low Signal-to-Noise Regimes

Multi-Objective Loss Formulation

Training deep financial models exclusively with Mean Squared Error forces predictions toward the unconditional empirical mean, dampening the variance needed to trade directional inflection points. Financial log-returns exhibit low signal-to-noise ratios and fat tails; an unconstrained quadratic penalty fits high-frequency gaussian noise while over-reacting to non-stationary outliers.

To maintain robust convergence, the MoE training objective uses a multi-task composite loss function that simultaneously balances magnitude precision, directional classification accuracy, and cross-sectional rank correlation:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{Huber}}(y_t, \hat{y}_t) + \lambda_{\text{dir}} \mathcal{L}_{\text{Directional}}(y_t, \hat{y}_t) - \lambda_{\text{IC}} \mathcal{L}_{\text{IC}}(y_t, \hat{y}_t) + \lambda_{\text{balance}} \mathcal{L}_{\text{Balance}}$$

The Huber loss provides robust regression against extreme price jumps by transitioning from quadratic to linear penalties when error magnitudes exceed threshold δ_H :

$$\mathcal{L}_{\text{Huber}}(e_t) = \begin{cases} \frac{1}{2} e_t^2, & \text{for } |e_t| \leq \Delta_H \\ \Delta_H \left(|e_t| - \frac{1}{2} \Delta_H \right), & \text{otherwise} \end{cases}$$

The directional penalty explicitly penalizes sign mismatches between forecasts and realized returns, discouraging trades executed against market momentum:

$$\mathcal{L}_{\text{Directional}}(y_t, \hat{y}_t) = \frac{1}{N} \sum_{t=1}^N \max\left(0, -\text{sign}(\text{span}_{42}(\text{start_span})[\text{span}_{42}(\text{end_span})[\text{span}_{49}(\text{start_span})[\text{span}_{49}(\text{end_span})](y_t) \cdot \hat{y}_t])\right)$$

To align predictions across multi-asset cross-sections, the Information Coefficient (Pearson correlation) loss directly maximizes the cross-sectional ranking accuracy of the forecasts:

$$\mathcal{L}_{\text{IC}} = \frac{\sum_{m=1}^M (y_{m,t} - \bar{y}_t)(\hat{y}_{m,t} - \bar{\hat{y}}_t)}{\sqrt{\sum_{m=1}^M (y_{m,t} - \bar{y}_t)^2} \sqrt{\sum_{m=1}^M (\hat{y}_{m,t} - \bar{\hat{y}}_t)^2}}$$

Router Regularization Without Pathological Balance

Standard MoE models often suffer from router collapse, where the gating network defaults to routing all tokens through a single dominant expert. In Natural Language Processing, this is mitigated by imposing an auxiliary loss that enforces uniform expert utilization. However, enforcing uniform routing in quantitative finance degrades performance: market crises and liquidity shocks occur far less frequently than steady-state drift regimes. Forcing an MoE to use tail-risk hedging experts equally during calm periods forces defensive strategies into benign markets, creating unnecessary portfolio churn and transaction costs.

To prevent router collapse while accommodating regime asymmetry, the framework applies a variance-matching auxiliary loss that regularizes routing variance without enforcing uniform utilization:

$$\mathcal{L}_{\text{Balance}} = E \sum_{i=1}^E m_i \cdot P_i, \quad \text{where } P_i = \frac{1}{T} \sum_{t=1}^T g_i(z_t), \quad m_i = \frac{1}{T} \sum_{t=1}^T \text{Softmax}_i(W_g z_t + \epsilon_t)$$

where $\epsilon_t \sim \mathcal{N}(0, \sigma^2)$ is exploration noise injected during the forward routing pass. This penalty ensures that every specialized expert receives sufficient gradient flow to remain active and calibrated, without artificially forcing crisis-oriented modules into regular rotation during low-volatility regimes.

Leakage-Proof Validation and Statistical Hypothesis Testing

Combinatorial Purged Cross-Validation

Standard K-Fold cross-validation assumes that samples are independent and identically distributed (i.i.d.). Applying this method to financial time series introduces substantial look-ahead bias and data leakage. Because forward-looking multi-step labels overlap across consecutive observations, training subsets inevitably contain labels whose evaluation horizons intersect with test partitions. Furthermore, serial correlation in feature space leaks information across train and test boundaries even without label overlap.

Standard walk-forward validation avoids look-ahead bias by testing sequentially through time,

but it evaluates only a single historical trajectory. This single-path dependency introduces high variance, rendering walk-forward validation susceptible to overfitting on the specific historical path observed.

Combinatorial Purged Cross-Validation (CPCV) resolves these issues by generating multiple out-of-sample backtest paths under strict independence guarantees. A historical timeline of T observations is partitioned into N contiguous, non-overlapping temporal blocks. For a specified test set size of k blocks, the algorithm constructs all

$\binom{N}{k}$ possible train-test combinations.

Within each combination, two data-cleansing operations are enforced:

- **Purging:** Training observations whose label formation windows overlap with any portion of the testing partition are removed from the training split. If a training label at timestamp t spans the interval $[t, t + h]$, and the test partition begins at t_{test} , any observation satisfying $t + h \geq t_{\text{test}}$ is purged.
- **Embargoing:** Autoregressive memory and market reaction lags can leak post-test structures back into subsequent training splits. An embargo removes a fixed percentage of observations (typically 1% to 5%) immediately following the end of each test window.

CPCV reconstructs $\varphi = \binom{N}{k}$

$\frac{k}{N}$ complete, distinct out-of-sample backtest paths. Rather than producing a single point estimate of performance, CPCV generates an empirical distribution of Sharpe ratios, maximum drawdowns, and information coefficients. This allows quantitative researchers to evaluate the probability of strategy failure across a wide range of alternative historical market sequences.

Controlling False Discoveries via the Deflated Sharpe Ratio

Evaluating multiple hyperparameters, model backbones, and routing configurations inflates the maximum observed Sharpe ratio through selection bias under multiple testing. A high Sharpe ratio identified via an unconstrained parameter sweep often reflects random noise rather than genuine market edge. The Deflated Sharpe Ratio (DSR) provides a statistical correction that tests whether an observed Sharpe ratio remains statistically significant after accounting for the number of trials, return non-normality, and sample length.

Under the null hypothesis that none of the evaluated strategy variations possess predictive skill, the expected maximum Sharpe ratio across N independent trials is approximated by:

$$E[\max_n \{\widehat{SR}_n\}] \approx \sqrt{2 \ln N} \left(1 - \frac{\gamma}{2 \ln N}\right) + \frac{\gamma}{\sqrt{2 \ln N}}$$

where $\gamma \approx 0.57721566$ represents the Euler-Mascheroni constant.

When parameter trials are correlated—as is typical in neural network architecture sweeps—treating nominal trials as independent overstates the multiple-testing penalty. The effective number of independent trials N_{eff} is derived from the eigenvalue spectrum of the strategy return correlation matrix C :

$$N_{\text{eff}} = \frac{\left(\sum_{k=1}^K \lambda_k\right)^2}{\sum_{k=1}^K \lambda_k^2} = \frac{K^2}{\text{Tr}(C^2)}$$

The Deflated Sharpe Ratio evaluates the Probabilistic Sharpe Ratio (PSR) against this multiple-testing benchmark $SR_0 =$

$$E[\varphi]$$

$\max(\widehat{SR})$:
 $\text{DSR} = \Phi\left(\frac{\widehat{SR}^* - SR_0}{\sqrt{\frac{\widehat{\gamma}_3}{T-1}}}\sqrt{1 - \widehat{\gamma}_3}\right) + \frac{\widehat{\gamma}_4 - 1}{4}\widehat{SR}^{*2}$
 where $\widehat{SR}^*$ is the highest observed annualized Sharpe ratio, T is the number of return observations, $\widehat{\gamma}_3$ is strategy return skewness, $\widehat{\gamma}_4$ is return kurtosis, and $\Phi(\cdot)$ denotes the standard normal cumulative distribution function. An MoE architecture must clear a significance threshold of $\text{DSR} \geq 0.95$ to reject the null hypothesis of selection bias, confirming that its observed performance reflects true predictive alpha.

Cross-Validation Protocol	Leakage Controls	Combinatorial Paths	Primary Statistical Output	Critical Vulnerability
Standard K-Fold	None (i.i.d. assumption)	None (single fold)	Single point-estimate metric	Severe look-ahead bias and inflated Sharpe estimates
Purged K-Fold	Label purging & post-test embargo	Single pseudo-path	Mean and variance across K folds	Limited sampling of alternative market regimes
Standard Walk-Forward	Expanding/rolling historical anchor	Single historical path	Single continuous equity curve	High path variance; overfits to historical chronology
Combinatorial Purged (CPCV)	Full purging & embargo across all $\binom{N}{k}$ splits	$\binom{N}{k}\frac{k}{N}$ complete paths	Empirical distribution of metrics & PBO	High computational overhead across large split sets

Production System Architecture, Drift Surveillance, and Execution Safeguards

Dual-Loop Production Architecture

Production financial deployment requires balancing sub-millisecond execution requirements against computationally intensive model updating. A dual-loop system architecture resolves this operational trade-off by decoupling real-time streaming inference from asynchronous model adaptation.

The online inference loop processes real-time tick feeds, order book updates, and macro releases through pre-compiled, optimized feature pipelines. The gating network generates dynamic soft routing allocations across the active expert ensemble, and the resulting predictions pass through pre-trade risk and volatility-scaling modules before being routed to execution venues.

Simultaneously, an asynchronous offline retraining loop operates out-of-band across historical data lakes. Realized execution fills, transaction costs, and tracking errors are streamed to an episodic memory buffer. The offline loop runs counterfactual credit assignment to update router query-key memory matrices, trains new modular policy vectors during persistent regime shifts, and validates candidate parameter sets using CPCV before promoting them to production.

Multi-Resolution Concept Drift Detection

Financial regimes shift across diverse temporal frequencies. Flash crashes and geopolitical escalations introduce sharp high-frequency shocks, whereas interest rate and commodity super-cycles unfold gradually over months. Single-scale statistical drift monitors either flood systems with false positives or respond too slowly to sudden market crises. Multi-Resolution Drift Detection (MRDD) monitors distribution shifts across multiple timeframes simultaneously using discrete wavelet transforms (DWT) and multi-scale Kolmogorov-Smirnov (KS) tests. The system tracks wavelet energy differences across decomposition scales $j \in \{1, \dots, J\}$:

$$D_{\text{wavelet}}^{(j)} = \frac{1}{2^j} \sum_k |d_{j,k}^{\text{test}}|^2 - \frac{1}{2^j} \sum_k |d_{j,k}^{\text{ref}}|^2$$

where $d_{j,k}$ represents the detail coefficients at scale j . Concurrently, a nonparametric Page-Hinkley cumulative monitor tracks model Information Coefficient (IC) decay:

$$U_t = \sum_{i=1}^t (\text{IC}_i - \overline{\text{IC}} + \delta_m), \quad m_t = \min_{1 \leq i \leq t} U_i$$

When the cumulative divergence $U_t - m_t$ exceeds an empirical control threshold $\lambda_{\text{threshold}}$, a concept drift alarm is triggered. This initiates an automated retraining cycle, signaling the offline loop to isolate a new policy delta vector for the emerging regime.

Deterministic Runtime Safeguards and Capital Preservation

To prevent execution failures caused by model hallucinations, router collapse, or flash volatility spikes, the production engine enforces deterministic runtime guardrails:

- Gating Entropy Fallback:** The Shannon entropy of the soft routing distribution is monitored continuously: $\mathcal{H}(g(z_t)) = -\sum_{i=1}^E g_i(z_t) \ln(g_i(z_t))$. If extreme market conditions cause the router distribution to collapse into a single expert ($\mathcal{H} < \mathcal{H}_{\text{critical}}$), the system triggers a fallback protocol, defaulting to an equal-weighted ensemble across all defensive modules to preserve predictive stability.
- Volatility-Targeted Sizing:** To protect against volatility clustering and liquidity contractions, position sizes are inversely scaled by prevailing realized volatility estimated via a GARCH(1,1) model: $\text{Target Position}_t = \text{sign}(\hat{y}_t) \cdot \min\left(\frac{\text{Cap}_{\text{max}}}{\sigma^* \hat{\sigma}_t^{\text{GARCH}}}, \dots\right)$ where σ^* is the institutional target volatility. This dynamically scales down portfolio exposure during high-volatility regimes.
- Drawdown Circuit Breakers:** A three-tiered trailing drawdown monitor tracks portfolio-level performance in real time. Crossing a 3% trailing drawdown reduces overall gross leverage by 50%; crossing a 5% drawdown forces complete liquidation to cash; and crossing an 8% drawdown shuts down live execution entirely, triggering an automated offline retraining and validation cycle.

Operational Subsystem	Core Mechanism	Target Latency	Trigger Threshold	Primary Failure Guarded Against
Real-Time	Soft MoE Routing	<10 ms	Continuous event	Execution latency

Operational Subsystem	Core Mechanism	Target Latency	Trigger Threshold	Primary Failure Guarded Against
Inference	with CAW		tick	slippage
Entropy Monitor	Shannon Gating Entropy $\mathcal{H}(g(z_t))$	Synchronous	$\mathcal{H} < \mathcal{H}_{\text{critical}}$	Gating network collapse into single expert
Drift Surveillance	Multi-Resolution Wavelet Energy	Daily / Hourly batch	$U_t - m_t > \lambda_{\text{threshold}}$	Undetected structural market regime shift
Volatility Sizing	GARCH(1,1) Inverse Scaling	Pre-trade check	$\hat{\sigma}_t > \text{Target } \sigma^*$	Capital destruction during volatility spikes
Circuit Breakers	Three-Tiered Trailing Drawdown	Real-time mark-to-market	3% / 5% / 8% MDD limits	Severe strategy breakdown and systemic tail events

Synthesis and Implementation Blueprint

Deploying a self-improving Mixture of Experts across equity indices and commodity futures requires addressing the challenges of market non-stationarity, low signal-to-noise ratios, and regime transitions at every stage of model design.

The system architecture combines: * **Information-Matched Decoupling:** Passing non-stationary macroeconomic and market volatility variables strictly through the routing gating network, while base predictors and localized residual experts operate exclusively on stationary, fractionally differentiated asset-level features.

- **Market-Specific Feature Engineering:** Capturing order flow imbalance on intraday intervals, adjusting for term structure roll yields and speculative positioning in commodities, and using CUSUM event-driven sampling to align feature extraction with market information arrival.
- **Continual Self-Improvement:** Training routing policies against the Differential Sharpe Ratio via reinforcement learning, preventing catastrophic forgetting through modular policy delta libraries, and updating router representations through counterfactual expert credit assignment.
- **Leakage-Proof Validation:** Enforcing combinatorial purged cross-validation with post-test embargoes to evaluate full distributions of out-of-sample paths, verified against selection bias using the Deflated Sharpe Ratio.
- **Production Guardrails:** Running a dual-loop infrastructure with multi-resolution drift detection, entropy fallback floors, and volatility-targeted risk controls.

Pipeline Stage	Architectural Decision	Implementation Methodology	Expected Performance Impact
1. Feature Ingestion	Fractional differencing (d^*) & CUSUM sampling	Expanding ADF grid search; dynamic volatility thresholds	Maximum memory retention under strict stationarity

Pipeline Stage	Architectural Decision	Implementation Methodology	Expected Performance Impact
2. Routing Architecture	RG-ResMoE with Correlation-Aware Weighting	Decoupled regime gating; cosine penalty on representations	Eliminates negative transfer and expert collinearity
3. Continual Adaptation	ReCAP policy library & counterfactual updates	Taylor series loss credit; isolated delta vectors	Eliminates catastrophic forgetting across regimes
4. Policy Optimization	Direct RL on Differential Sharpe Ratio	Recursive moment tracking with transaction costs	Direct optimization of risk-adjusted returns
5. Model Validation	CPCV with Deflated Sharpe Ratio	Purging label overlaps; N_{eff} multiple-testing test	Guarantees out-of-sample statistical significance
6. Production Surveillance	Dual-loop execution & Wavelet MRDD	Real-time CUSUM; multi-scale KS tests; circuit breakers	Sub-10ms inference with automated model retraining

By integrating specialized modular capacity, continuous parameter-isolated adaptation, and rigorous backtesting safeguards, this architecture provides a robust framework for systematic forecasting in institutional equity and commodity markets.

Works cited

1. RAVEN: A Regime-Aware Variable-context Expert Network ... - arXiv, <https://arxiv.org/html/2606.24062> 2. RefineBridge: Generative Bridge Models improve financial ... - arXiv, <https://arxiv.org/html/2512.21572v2> 3. Do Technical Indicators Improve Deep Learning Forecasts? An, <https://www.preprints.org/manuscript/202602.0692> 4. A Drift-Aware Dynamic Mixture of Experts Framework for Non ... - arXiv, <https://arxiv.org/html/2605.20678v1> 5. Continual Deep Learning for Time Series Modeling - MDPI, <https://www.mdpi.com/1424-8220/23/16/7167> 6. TradingMoE: Routing the Right Experts in Evolving Markets - arXiv, <https://arxiv.org/html/2608.11785> 7. Regime-Gated Residual Mixture-of-Experts for Cross-Sectional, <https://arxiv.org/html/2608.12251v1> 8. Reinforcement Learning for Trading - NIPS, <http://papers.neurips.cc/paper/1551-reinforcement-learning-for-trading.pdf> 9. Vectorbt Walk-Forward: Implementing Purged K-Fold Cross-Validation, <https://trader-algoritmico.com/blog/vectorbt-walk-forward-implementing-purged-k-fold-cross-validation> 10. Deflated Sharpe Ratio: Correcting for Multiple Testing, <https://paperswithbacktest.com/course/deflated-sharpe-ratio> 11. A Regime-Aware Variable-context Expert Network for Financial Time, <https://arxiv.org/abs/2606.24062> 12. MAMBAFORMER-MOE: MAMBA-BASED MIXTURE- - OpenReview, <https://openreview.net/pdf?id=55Ruu7bF4b> 13. Mamba for Time Series Analysis: A Contemporary Survey, <https://www.preprints.org/manuscript/202605.0995> 14. MambaFormer-MOE: Mamba-Transformer-based Mixture-of-Experts, <https://openreview.net/forum?id=55Ruu7bF4b> 15. Billion-Scale Time Series Foundation Models with Mixture of Experts, https://www.researchgate.net/publication/384296045_Time-MoE_Billion-Scale_Time_Series_Foundation_Models_with_Mixture_of_Experts 16. Mamba or Transformer for Time Series

Forecasting? Mixture ... - arXiv, <https://arxiv.org/html/2408.15997v1> 17. Multi-Scale Hybrid Mamba-Transformer Experts for Time Series, <https://www.semanticscholar.org/paper/SST%3A-Multi-Scale-Hybrid-Mamba-Transformer-Experts-Xu-Chen/4ea7378fc3cb0c6e7092ae162cbd1537ae637597> 18. Meta-Labeling in Quantitative Strategies | PDF - Scribd, <https://www.scribd.com/document/720604417/Does-Meta-Labeling-Add-to-Signal-Efficacy> 19. Machine Learning Trading Essentials (Part 2): Fractionally, <https://hudsonthames.org/machine-learning-trading-essentials-part-2-fractionally-differentiated-features-filtering-and-labelling/> 20. Free Order Flow Imbalance Calculator | OFI & Market ... - Pineify, <https://pineify.app/free-tools/order-flow-imbalance-calculator> 21. Order-Flow Imbalance (OFI) - Emergent Mind, <https://www.emergentmind.com/topics/order-flow-imbalance-ofi-7dff1686-44cf-4cf4-a602-b24df2b7c56e> 22. What is the order flow imbalance?, <https://quant.stackexchange.com/questions/43751/what-is-the-order-flow-imbalance> 23. Understanding Order Flow Toxicity - Medium, <https://medium.com/@simomenaldo/understanding-order-flow-toxicity-58fa317b3d01> 24. (PDF) Forecasting High Frequency Order Flow Imbalance, https://www.researchgate.net/publication/382944327_Forecasting_High_Frequency_Order_Flow_Imbalance 25. The Promise of Time-Series Foundation Models for Agricultural, <https://ideas.repec.org/p/arx/papers/2601.06371.html> 26. Tim Bollerslev - IDEAS/RePEc, <https://ideas.repec.org/e/c/pbo66.html> 27. The Deflated Sharpe Ratio: Correcting for Selection Bias, https://www.researchgate.net/publication/286121118_The_Deflated_Sharpe_RatioCorrecting_for_Selection_Bias_BacktestOverfitting_and_Non-Normality 28. Full article: Agricultural commodity futures prices prediction via long, <https://www.tandfonline.com/doi/full/10.1080/15140326.2019.1668664> 29. Regime-Adaptive Continual Learning for Portfolio Management, <https://www.alphaxiv.org/abs/2606.00143> 30. Learning to trade via direct reinforcement - ResearchGate, https://www.researchgate.net/publication/5602298_Learning_to_trade_via_direct_reinforcement 31. Optimizing Crypto-Trading Performance: A Comparative Analysis of, <https://www.mdpi.com/2227-7390/14/5/794> 32. Deep Reinforcement Learning in Factor Investment - arXiv, <https://arxiv.org/html/2509.16206> 33. Approximating Sharpe and Sortino ratios from Exponential moving, <https://quant.stackexchange.com/questions/57568/approximating-sharpe-and-sortino-ratios-from-exponential-moving-averages> 34. Does catastrophic forgetting negatively affect financial predictions?, <https://openreview.net/pdf?id=cvPCsagdRL> 35. Randomized Neural Network with Adaptive Forward Regularization, <https://arxiv.org/html/2510.21367v1> 36. ACLEGR: Adaptive Continual Learning with Efficient Generative, <https://www.techrxiv.org/doi/pdf/10.36227/techrxiv.175417352.25931330/v1> 37. Regime-Adaptive Continual Learning for Portfolio Management - arXiv, <https://arxiv.org/html/2606.00143v1> 38. Catchup results for Machine Learning on Fri, 29 May 2026 - arXiv, <https://arxiv.org/catchup/cs.LG/2026-05-29?abs=True&page=1> 39. Agentic Trading: When LLM Agents Meet Financial Markets - arXiv, <https://arxiv.org/html/2605.19337v1> 40. at CLEF 2026 - arXiv, <https://arxiv.org/html/2607.12233v1> 41. Continual Learning with Deep Generative Replay | Request PDF, https://www.researchgate.net/publication/317101000_Continual_Learning_with_Deep_Generative_Replay 42. purged-cross-validation/paper/paper.md at main - GitHub, <https://github.com/eslazarev/purged-cross-validation/blob/main/paper/paper.md> 43.

Combinatorial Purged Cross-Validation Insights | PDF - Scribd,
<https://www.scribd.com/document/725401650/SSRN-id4778909> 44. Vectorbt Walk-Forward with
Purged K-Fold: Avoiding Data Leakage,
<https://trader-algoritmico.com/blog/vectorbt-walk-forward-with-purged-k-fold-avoiding-data-leakage>
45. The Deflated Sharpe Ratio: correcting for how many tries you took,
<https://quanterlab.com/articles/qlway-deflated-sharpe> 46. Deflated Sharpe Ratio: Formula,
Intuition, and a Worked Example, <https://quanterlab.com/articles/foundations-dsr> 47. THE
DEFLATED SHARPE RATIO: CORRECTING FOR SELECTION,
<https://www.davidhbailey.com/dhbpapers/deflated-sharpe.pdf> 48. Deflated Sharpe Ratio:
Correcting for Backtest Overfitting, <https://quantmedia.io/learn/deflated-sharpe-ratio.html>